

# CIS-505

## Software Systems

Instructor: Matt Blaze

# Important Administrative Stuff

- This is CIS-505- Software Systems
  - We meet Tuesday & Thursday, 1330-1500
- There's a waiting list to enroll
  - contact Mike Felker with:
    - name, status (CIS PhD, CIS MS, other)
  - We'll let you know before next class
- This is a PhD WPE-I course
  - PhD students should be in the appropriate section

# Meet your staff...

- Your Instructor: Matt Blaze
  - Office: Levine 611
  - Office Hours: Tuesday 4:30 – 5:30 (pm!) & by appointment (email me!)
  - [blaze@cis.upenn.edu](mailto:blaze@cis.upenn.edu)
- That's me – I'm the person talking right now

# Textbooks, course web page

- A.S. Tanenbaum & Maarten van Steen.  
*Distributed Systems (2<sup>nd</sup> Edition)*. Prentice-Hall, 2007.
  - it's in the book store
- Optional: W. Richard Stevens and Stephen A. Rago.  
*Advanced Programming in the UNIX Environment (2/e)*. Addison-Wesley Professional. 2005
- Course web page is:  
`www.crypto.com/courses/fall12/cis505/`
  - Check frequently for updates and news
  - Write this URL down. Do it RIGHT NOW.

# Prerequisites

- Basic understanding of computer operating systems & architecture (CIS240, CIS380)
- Working knowledge of C language and Unix shell and programming environment
- Shell login on “speclab” Linux cluster
  - Unix shell account
  - Check to make sure you have a working account. (Do this TODAY – really)
- Read the above point and actually do it

# Grading

- Midterm (15%), Final (35%)
- Two *Substantial* programming projects (15% and 25%)
  - Done in *two-person* teams
  - putting things off until the last minute doesn't work
- Homeworks (10%)
- fascist, inflexible late homework policy strictly enforced!

# A quick word about “PowerPoint”

- Some courses are based on *slideware*
  - well defined set of PowerPoint slides in which the content of the course is contained
  - slides more important than textbook, lecture, readings, or homework
- This is not one of those courses
  - Slides useful here as outline of topics, at most
- Slides will be made available *after* each class
  - do not use as substitute


# Policy for collaboration

- Individual written homeworks:
  - **YES:** discuss with friends, acknowledge them, write your own solutions
  - **NO:** copy past-year solutions word-for-word (or with minor modifications)
  - **NO:** copy answers scribbled by someone on used textbook
  - **NO:** use search engines to look for past-year solutions
- Group projects
  - **YES:** discuss concepts and programming approaches on newsgroup
  - **NO:** modifying/submitting solutions from other groups current/past year
- All students must contribute fair-share to team projects:
  - Individually fill in a form listing contributions of each group member.
  - Logs in version control to resolve disputes



# This is a Distributed Systems Course!

- **Fundamentals:** how distributed systems are designed
  - Naming and addressing
  - Clock synchronization, logical clocks, events
  - Elections, distributed commits/recovery
  - Failures, Byzantine fault tolerance
- **Real-world insights:** how are real systems built?
  - How does theory translate into practice?
  - Network File System, Coda FS, Chord FS
  - “Backends” of search technologies
- **Skills:** distributed system implementation
  - Hands-on programming experience
  - Code repository (Subversion or CVS)
  - Low-level C programming to Java-based distributed programs



# First half of course:

## Fundamentals

- Concurrency: Processes, semaphores, monitors (review of basic OS concepts)
- Inter-process protocols, RPC, RMI, sockets, msg-queues
- Naming, addresses, locations
- Clock synchronization, logical clocks, clock vector
- Coordination and agreements: elections, mutual exclusion, quorums/voting, distributed commit, total ordering
- Consensus in the presence of faults / failures
- Reliable group communication
- Distributed snapshots / checkpoint
- Replication and consistency models

# Second half of course:

## Putting Theory into Practice

- Topics may depend on time and interests.
- Examples:
  - Distributed mail server (not covered in class, learn via project 2)
  - Distributed file systems (Network File System, Andrew File System, Coda File System, Chord File System, Ivy File System)
  - Content distribution networks (Akamai)
  - Large-scale clusters:
    - Google: MapReduce, Google File System, BigTable, Chubby distributed lock system
    - Amazon: Dynamo storage system
    - Microsoft: Dryad, DryadLINQ
    - (If time permits) Cloud computing: Amazon EC2 cloud, Hadoop middleware, consistency models
- Guest lectures on advanced topics

So... let's get started

Operating Systems:  
Isolating Programs

# Mostly, this course will be about how systems *communicate*

- But first, we need to answer some questions:
  - how does software run on computers?
  - how come you usually don't have to worry about other running programs?
- Operating systems *isolate* running programs
  - allow transparent timesharing
  - provide narrow communication channel
  - how does this work?

# Processes / Time Sharing

- OS allows program to act as if it “owns” the machine
  - not worry about other running programs
  - not use more than its fair share of resources
  - get services from the OS
  - communicate with others
- Exploits two hardware / CPU mechanisms:
  - relocatable memory
  - “kernel mode” and “user mode” CPU state
    - “traps”

# Relocatable Memory

- CPU feature that allows programs to not know in advance where they will be in memory
  - and also avoid stepping on each other's memory by mistake
- Many ways to this is implemented:
  - offsets, prefixes, virtual memory
  - details don't matter yet
- Every running program can have its own private *address space*

# “Supervisor mode” / “User mode” and “traps”

- When the CPU is in “supervisor” mode, it can do anything, and can address any part of memory.
  - the OS kernel runs in supervisor mode
  - supervisor mode can switch to user mode at will
- When the CPU is in “user” mode, it has access only to its own address space, and can’t talk directly to devices
  - User process run in user mode
  - But after a trap we can switch to supervisor mode...



# “Traps”

- A user mode program can switch to supervisor mode by issuing a “trap”
  - but there’s a catch... when it issues a trap it starts running OS code in the supervisor address space
  - user program can’t overwrite this code
- Traps are used to implement system calls to provide services to user processes that require:
  - communication with a hardware device
  - communication with another process or with the OS
- Traps also happen when devices issue *interrupts*
  - real time clock ticks, message arrives on network, etc.

# How an OS runs processes...

## (an oversimplification)

- OS must keep track of which process are assigned which sections of memory plus other stuff
- To run a new process, assign it some memory and put its code there
  - switch to user mode and start running at the first address of the program
- The OS keeps a record of every process
  - assigned memory is, current program counter, etc.
  - this is called the process' *context*
  - enough information to restart process where it left off

# Eventually, a trap happens

- Either because the running program issued it or because a device did (e.g., the clock ticked)
- First the OS records the state of the running process' context in its record
- Next it will do whatever servicing the trap requires
  - e.g., send something to the printer
- finally, it will pick a user process to restart
  - maybe the one that was running, maybe not
  - restarts based on the new process' context record
  - back in user mode now

# Processes and System Calls

- First we'll look at processes in (Unix-like Oss) from the *user's* perspective
  - what is a process?
  - what's the interface to the OS for managing them?
- Then we'll look at how all this is implemented by the OS

# What's a process?

- A “*program in execution*”
  - with associated (data and execution) context
- *Not* the same as “program” or “application”
  - a given program may be running 0, 1, or >1 times
- Each *instance* of a program is a separate process
  - with its own address space and other context
- Some OSs and GUIs obscure this distinction from the user. Don't be fooled.
  - e.g., what happens when you double click the browser icon the first time? the second time?

# Unix Process Hierarchy

- A Unix process comes in to being when another process creates it
  - a newly created process is the “child” of the “parent” that created it
  - every process has exactly one immediate parent
  - a process might create any number of children
- “Root” process, *init*, created when OS is booted
  - every process can be traced back to *init*
- Processes have a unique *process ID* (PID) number
  - index in the OS to the processes’ context

# Creating a process: fork()

- The fork() system call creates a new process
- The child process is an (almost) exact clone of the parent (with it's own copy of parent's address space)
  - starts running as soon as fork() returns
  - both child and parent now running simultaneously
- What good is this?
  - write code to *behave differently* if you're the child
- How can you tell if you're the child or the parent?
  - for the child, fork() return value is 0
  - for the parent, fork() return value is the PID of child

# Replacing a process: `exec()` and friends

- The `exec()` system call (and variants) *replaces* a process with a new program
  - it doesn't create any new processes
  - the new program is specified by the name of the file containing its executable code plus arguments
- The old code stops running as soon as it calls `exec()` (if the executable file is successfully run)
- What good is this?
  - usually run after `fork()`



# Creating a child process

- Typically uses both `fork()` and `exec()`

```
pid=fork() ;
```

```
if (pid != 0)
```

```
    /* do parent stuff */
```

```
else
```

```
    exec("/bin/child") ;
```

# Annoying details

- Exec is really a complicated family of system calls, with slightly different versions of the way each processes arguments
  - you specify the name of the executable file and its arguments (in an array of string pointers)
  - see the manual pages for `execl()` and `execve()`
- Fork might fail
- Exec might fail

# Wait()

- Often a process will have nothing to do until its child terminates
  - e.g., what the shell usually does when you type a new command
- `wait()` *blocks* until the child terminates
  - so it doesn't return immediately
- Annoying details: status argument, what if there's no child, etc.

# Creating a child process (slightly refined version)

- Using fork(), execve() and wait()...

```
pid=fork() ;
```

```
if (pid != 0)
```

```
    wait(&status) ;
```

```
else
```


```
    if (execve("/bin/child"... ) !=0)
```

```
        exit(1) ;
```

# Wait() just a second here: What was this about “blocking?”

- Some system calls can “return” immediately
  - fork() and exec(), for example
- Others are asking for something that takes a while
  - wait(), for example
  - no sense in having the OS restart the process until the request is ready
- A process might be in one of three basic states:
  - *running*, *blocked* and *ready*
  - more on this later

# How the Unix shell runs commands



- when you type a Unix “command” like “ls”, your shell process first *forks* an (almost) identical clone of itself
- the new “child” process makes an *exec* call (to the ls program), which causes it to stop executing the shell and start executing your command
- the original “parent” process, still running the original shell, *waits* for the child process to terminate

# Communicating with processes: `exit()`, `kill()`

- `exit()` forces the current process to terminate
  - takes a return value that can be communicated back to parent via `wait()`
  - also called automatically when process ends
- `kill()` sends a *signal* to a process
  - if the process has set up a *signal handler*, it's called, much like an interrupt
    - otherwise the process terminates
  - restrictions on sending signals
    - only can send signals to the same user's processes

# More system calls:

## I/O with read() and write()

- Much I/O is based on a streaming model
  - sequence of bytes
- write() sends a stream of bytes somewhere
- read() blocks until a stream of input is ready
- Annoying details:
  - might fail, can block for a while
  - file descriptors...
  - arguments are pointers to character buffers
  - see the read() and write() man pages



# File descriptors

## open(), close()

- A process might have several different I/O streams in use at any given time
- These are specified by a kernel data structure called a *file descriptor*
  - each process has its own table of file descriptors
- Open() associates a file descriptor with a file
- Close() destroys a file descriptor
- Standard input and standard output are usually associated with a *terminal*
  - more on that later

# More I/O stuff

- It's possible to hook the output of one program into the input of another
  - pipe()
- It's possible to block until one of several file descriptor streams is ready
  - select()
- Special calls for dealing with network
  - sockets, etc.

# What's the point here?

- System calls are the main interface between processes and the OS
  - system calls are like an extended “instruction set” for user programs that hide many details
  - first Unix system had a couple dozen system calls
  - current systems have a couple *hundred*
- Understanding the system call interface of a given OS lets you write useful programs under it
- Natural questions to ask:
  - is this the right interface? how to evaluate?
  - how can these system calls be implemented?

# Implementing processes in Unix

- The OS kernel manages processes
- Has to solve various problems, including:
  - keeping track of the processes' states
    - including enough information to stop and restart them
  - keeping track of the available resources
    - those allocated to specific processes
    - reclaiming resources when a process releases them (or exits)
  - deciding who to run next (scheduling/dispatching)
  - managing communication
    - pipes, file I/O, etc.

# Keeping track of process state

- A kernel data structure, usually called the *process table*, keeps track of each active process
  - entries in the table are *process control blocks* (PCBs), which contain data on state of each process
  - some of the data is directly in the PCB (e.g., registers)
  - other data in form of pointers (e.g., memory)
- In other words, a process' *context* is described by its process control block
  - this allows the OS kernel to stop and restart processes from where they left off


# Process Context

- The full context of a process includes:
  - contents of main memory
  - contents of CPU registers, including the program counter, etc.
  - other info (open files, I/O in progress, etc.)
- Main memory -- three logically distinct regions of memory:
  - text region: contains executable code (typically read-only)
  - data region: storage area for dynamically allocated data structure, e.g., lists, trees (typically heap data structure)
  - stack region: run-time stack of activation records
- Registers: general registers, PC, SP, PSW, segment registers
- Other information:
  - open files table, status of ongoing I/O
  - process status (running, ready, blocked), user id, ...

# Who's running when?

- Most of the time, a regular (user mode) process is running
- When a system call or an interrupt occurs, the kernel runs (in kernel mode)
  - as soon as it's finished, the *scheduler / dispatcher* picks a new process to run
  - it uses the PCB of that process to restart it

# Basic Process States



1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available

- Possible process states
  - Running: executing (only one per CPU)
  - Blocked: waiting (e.g., for I/O)
  - Ready: waiting to be scheduled



# Implementing system calls

- When a system call or interrupt occurs:
  - first do whatever is requested
    - schedule I/O, update real time clock, etc.
  - maybe move some processes from *blocked* to *ready*
  - go through the list of *ready* processes, pick one and restart it based on its PCB (in the process table)
    - how to pick? it's a policy question.
      - fairness, efficiency, etc.
      - we'll discuss in detail later

# Implementing Unix System Calls (oversimplified)

- `fork()`
  - duplicate process table entry, allocate and copy memory
  - both parent and child become ready immediately
- `exec()`
  - open file, load (or map) contents to memory, reset context
  - process goes to ready state as soon as this happens
  - might block if opening or loading the file takes a while

# Implementing Unix System Calls (continued)

- `wait()`
  - if child not running, put process in ready state
  - if child is running, put process in blocked state
- `exit()`
  - terminate process and release resources
  - if parent is blocked in `wait()`, change parent's state to ready

# Implementing Unix System Calls (continued)

- `read()`
  - if error, put process in ready state
  - if input is available, put process in ready state
  - if no input available, put process blocked state
- `write()`
  - if error, go to ready state
  - if I/O channel not available, put process in blocked state, otherwise put process in ready state
  - if another process was blocked waiting for I/O, put that process in ready state
- On many OSs, I/O goes through a *scheduler*

# Other system calls

- `signal()`, `sleep()`, `select()`, etc all ask OS kernel to perform specific functions for process
  - block until something happens
  - send a message to another process
  - interact with I/O devices
  - etc
- System calls documented in section 2 of the Unix / Linux manual
  - type “man 2 read” for example

# A question...

What's the difference between a system call (like `write()`) and a standard library function (like `printf()`)?

# What's the point here?

- System calls are the main interface between processes and the OS
  - system calls are like an extended “instruction set” for user programs that hide many details
  - first Unix system had a couple dozen system calls
  - current systems have a couple *hundred*
- Understanding the system call interface of a given OS lets you write useful programs under it
- Natural questions to ask:
  - is this the right interface? how to evaluate?
  - how can these system calls be implemented?

# So far...

- System calls are the interface that processes (running programs) use to communicate with the OS
  - called like functions, but implemented in the OS
- We've looked at the basic Unix process management system calls
  - `fork()`, `exec()`, `wait()`, `kill()`, `exit()`
- And some I/O system calls
  - `read()`, `write()`



# Why bother?

- Understanding the system call interface of a given OS lets you write useful programs under it
- Natural questions to ask:
  - is this the right interface? how to evaluate?
  - how can these system calls be implemented?

# Implementing processes in Unix

- The OS kernel manages processes
- Has to solve various problems, including:
  - keeping track of the processes' states
    - including enough information to stop and restart them
  - keeping track of the available resources
    - those allocated to specific processes
    - reclaiming resources when a process releases them (or exits)
  - deciding who to run next (scheduling/dispatching)
  - managing communication
    - pipes, file I/O, etc.

# Keeping track of process state

- A kernel data structure, usually called the *process table*, keeps track of each active process
  - entries in the table are called *process control blocks* (PCBs), which describe the running state of each process
  - some of this is stored in the PCB (e.g., registers)
  - other data in form of pointers (e.g., memory)
- In other words, a process' *context* is described by its process control block
  - this is enough information for the OS kernel to stop and restart processes from where they left off

# Process Context

- Process' full context (process table entry) includes:
  - the memory it's using for code and data
  - contents of CPU registers, including the program counter, etc.
  - other things (open files, I/O in progress, etc.)
  - details depend on the particular hardware architecture
- Memory -- three logically distinct regions of memory:
  - text region: contains executable code (typically read-only)
  - data region: storage area for dynamically allocated data structure, e.g., lists, trees (typically heap data structure)
  - stack region: run-time stack of activation records
- Registers: general registers, PC, SP, PSW, etc.
- Other things:
  - open files table, status of ongoing I/O
  - process status (running, ready, blocked), user id, ...

# Running Programs: Who's running when?

- Most of the time, a (*user mode*) process is running
- When a system call or an interrupt occurs, the kernel runs (in *kernel mode*)
  - kernel does what it has to, then the *scheduler / dispatcher* part of the kernel picks the next process to run
    - maybe the same one, maybe a different one
  - scheduler uses the PCB (process table entry) of the selected process to find out how to restart it

# Non-Blocking vs. Blocking

- Some system calls are “non-blocking”
  - OS can return to the calling process immediately
  - ask the OS to do something, but start running again right away
  - `fork()`
- But others explicitly “block” the calling process until something happens
  - process goes to sleep until conditions have changed to where it makes sense to run
  - `wait()`, `read()`

# Blocking let's us avoid “polling” system calls

```
pid=fork();  
if (pid != 0) /* parent */  
    wait(&status); /* parent blocks  
    */  
else /* child */  
    if (execve("/bin/child"... ) !=0)  
        exit(1);
```

Why is this a good way to do things?  
(hint: think about *The Simpsons*)


# What if wait() didn't block?

- We'd need an "ischilddeadyet()" syscall
  - returns TRUE if child is dead, else FALSE

```
pid=fork();  
if (pid != 0) /* parent */  
    while (ischilddeadyet(pid)==FALSE)  
        ; /* "are we there yet?" */  
else /* child */  
    if (execve("/bin/child"...)!=0)  
        exit(1);
```



# Basic Process States



1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available

- Processes can be in any of three states:
  - **Running**: currently executing; at most one per CPU
    - may become *blocked* or *ready* after a system call
  - **Blocked**: waiting for something (e.g., I/O)
    - currently *not* eligible to be scheduled to run
  - **Ready**: not waiting for anything
    - not running, but eligible to be scheduled to run

# What does the OS do when a system call occurs?

- First, do whatever actions are requested
  - schedule I/O, update real time clock, etc.
  - based on these actions, we might move some processes from *blocked* to *ready*
- Next, go through the list of *ready* processes and select one to run next
  - how to pick? different OSs have different policies.
    - fairness, efficiency, etc.
    - we'll discuss in detail later
- Finally, restart selected process based on its PCB entry (in the process table)

# Implementing Unix System Calls: Process Management

- `fork()`
  - duplicate process table entry, allocate and copy memory
  - both parent and child become *ready* immediately
- `exec()`
  - open file, load (or map) contents to memory, reset context (e.g., set PC to first instruction)
  - process usually goes to *ready* state as soon as this happens
    - might *block* if opening or loading the file takes a while

# Implementing Unix System Calls: Process Management cont'd

- `wait()`
  - if child not running, put process in *ready* state
  - if child is running, put process in *blocked* state
- `exit()`
  - terminate process and release resources
    - remove PCB entry
  - if parent is blocked in `wait()`, change parent's state to *ready*

# Implementing Unix System Calls: I/O

- `read()`
  - if error, put process in *ready* state
    - return error code
  - if input is available, put process in *ready* state
  - if no input available, put process *blocked* state
- `write()`
  - if error, go to *ready* state
  - if I/O channel not available, put process in blocked state, otherwise put process in *ready* state
  - if another process was blocked waiting for input from this write, put that process in *ready* state
- On many OSs, I/O doesn't happen immediately
  - goes through a *scheduler*

# Other system calls

- `signal()`, `sleep()`, `select()`, etc ask OS kernel to perform various other services for process
  - *block* until something happens
  - send a message to another process
  - interact with I/O devices
  - etc
- System calls are documented in section 2 of the online Unix / Linux manual
  - try “**man 2 read**” for example

# Details: Interrupts, System Calls and Processes

- A more detailed look at how basic OS services are implemented on real CPUs
  - Traps and system calls
  - User/supervisor mode
  - Context switching
  - Serving interrupts
  - Managing processes
- Efficiency issues

# What's a system call, really?

- Looks a lot like a function as far as the programmer is concerned
  - parameters, return value, etc
  - function call wrapper
- Main difference is that the system call executes code in the OS (in supervisor mode)
  - In this sense, can think of system call as like an extension of the instruction set
- How to communicate parameters and return value between the program and the OS?



# Calling a regular function (not a system call)

- Call (typical CPU architecture):
  - Push calling data onto stack
    - parameters
    - state (general registers, control registers, etc)
  - Jump to (load PC with address of) first instruction of function
- Return:
  - pop state off stack and reload registers
  - push return value onto stack
  - jump to next instruction of caller (as read from stack)

# Invoking a system call is similar, except:

- The code for the “function” is in the OS
  - operates in **supervisor** mode
  - different address space
- Limited number of entry points (“functions”)
  - a user process can’t just jump to anywhere it likes in the OS
  - enforced by semantics of TRAP instruction
- A system call might not return right away
  - process’ state might change to *blocked* as a result of the system call’s actions
  - OS might decide to schedule a different program

# Calling a system call

- Push data onto stack
  - parameters to system call
  - system call number
  - registers, current state, etc
- Execute a TRAP instruction
  - this causes several things to happen
    - change CPU mode bit from **user** to **supervisor**
    - jump to specific address in OS address space
      - as indexed by system call number

# Calling a System Call (continued)

- OS pops and saves calling process's state from user stack
  - saved in process table entry
- OS pops syscall number and parameters (from stack)
- OS calls appropriate system call function
  - pushes return value onto process's stack
- OS calls scheduler/dispatcher to select and load next process to run
  - maybe the one that issued the syscall, maybe not

# Returning from a system call

- OS sets mode bit back to USER
  - and reloads state registers (PC, etc)
- Now we're back in the calling process
  - next instruction is the one after the TRAP that issued the syscall
- Process pops syscall return value from stack and gets on with its business

# Next class

- Concurrency:
  - Processes, threads, semaphores, monitors
  - First homework assigned
- Readings (which you should have done already)
  - Chapters 1-3
- Your TODO list:
  - Log onto speclab cluster
  - Get your textbook!
  - Visit course website: <http://www.crypto.com/courses/spring12/cis505/>
  - Note Homework 1:
    - <http://www.crypto.com/courses/spring12/cis505/hw1.html>
  - Start looking for a project partner.