

CIS-505 – Software Systems

Notes for 1/19/12

© Matt Blaze

A quick reminder

- Homework 1 due Wednesday 1/25, 11:59pm
 - Yikes! Less than a week!
 - This is an *individual* assignment (not a group project)
- If you're going to run experiments, please don't use the Eniac cluster
 - use the “speclab” cluster for dangerous stuff!

Processes and Threads

- What's a process?
 - “a program in execution”
 - a memory address space (containing code & data)
 - various other resources (e.g. open files)
 - state information (pgm counter, registers, stack pointer etc)
 - the stuff stored in the process' PCB
 - really *two* categories of things
 - a *collection of resources*
 - the code & address space, open files, etc.
 - a *thread of execution*
 - the current state that operates on these resources
- We can think about these two things separately

Threads in the process model

- A lot of what the OS does is intended to keep processes from interfering with each other
 - every thread of execution is associated with its own grouping of resources (process)
 - I can't write over your process's address space
- This is nice, but it means that if you want to change threads, you must switch processes
 - requires OS intervention & expensive context switch
 - this is a shame, because some applications logically consist of more than one thread but need only one grouping of resources

Multi-threaded applications

- Could we let multiple threads share a common memory address space?
 - for applications with
 - a need to share data structures among threads
 - no need to for the OS to enforce resource separation (because the threads “trust” each other)
 - *not* for arbitrary code / general programs
- Some potentially *multi-threaded* applications:
 - web *server*
 - serves pages to several different clients at once
 - web *browser*
 - load different pages simultaneously

Can we implement multiple threads in a single process?

- We can do in user mode many of the functions usually handled by the OS
 - assumes the threads are cooperating so we don't need hardware enforcement of separation
- Basic idea: a “dispatcher” subroutine (in the process) that is called when a thread is ready to relinquish control to another thread
 - manages stack pointer, program counter
 - can switch process's internal state among threads

Inter-process communication

- The process model is a useful way to *isolate* running programs
 - separate resources, state, etc
 - narrow communication channel (wait, kill, etc)
 - vastly simplifies most programs - no need to worry about what other processes are doing
- Unfortunately, some applications work best if multiple threads are allowed to more tightly communicate and synchronize with each other
 - and this can make things complicated again

When might threads need to communicate?

- Many problems all over operating systems
 - threads with access to same data structures
 - kernel/OS access to user process data
 - processes sharing data via shared memory
 - processes sharing data via system calls
 - processes sharing data via file system
- ...and computer science generally
 - database transactions
 - programming languages that support parallelism

What makes this hard?

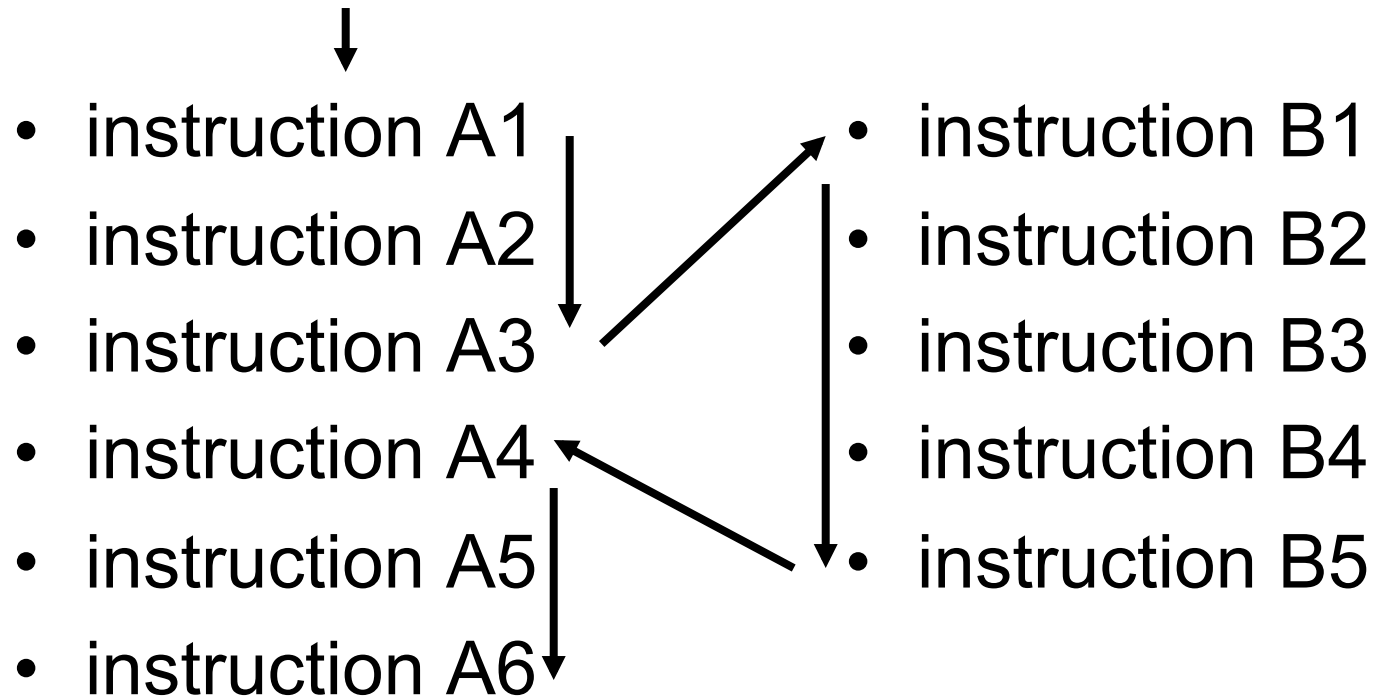
- Process model allows more than one program to run “at the same time”
 - they’re not *actually* running at the same time (on uniprocessors), of course
 - “at the same time” means an arbitrary -- and unpredictable -- interleaving of the machine instructions of the “simultaneous” processes
- Problem: logical operations on shared data often involve more than one instruction
 - your process might be stopped (by the OS) and another process might run and alter shared data you were in the middle of operating on

Two “simultaneous” processes

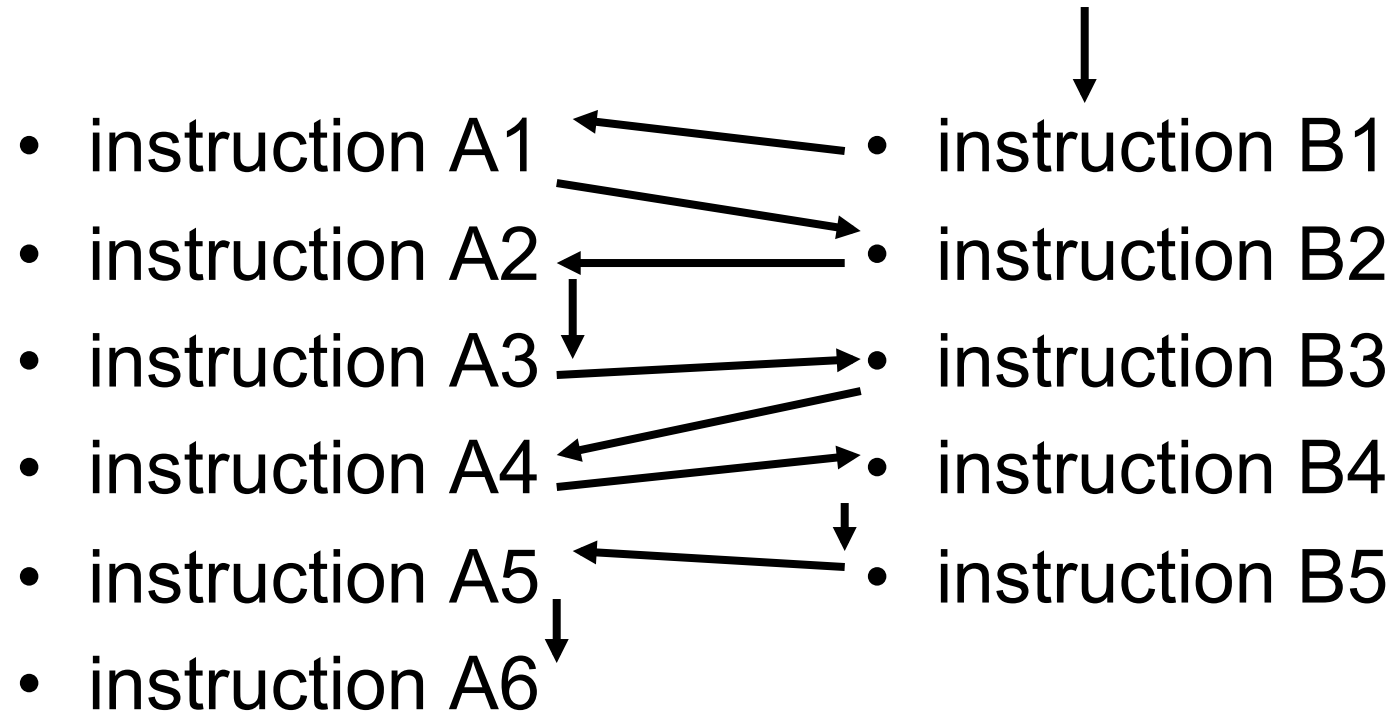
A (A1-A6) and B (B1-B5)

- instruction A1
 - instruction A2
 - instruction A3
 - instruction A4
 - instruction A5
 - instruction A6
- instruction B1
 - instruction B2
 - instruction B3
 - instruction B4
 - instruction B5

...could be scheduled like this:



... or maybe like this:



Normally, this wouldn't bother us

- The two threads are isolated, after all
 - different memory, registers, etc
 - can't distinguish between different scheduling sequences, so no problem
- But what if the two threads share access to the same memory?
 - this is where the trouble begins...

Race conditions and Synchronization

- Operations on shared data structures often consist of short “bursts” of instructions
- When two processes/threads are executing concurrently, the result can depend on the precise interleaving of the two instruction streams (this is called a **race condition**)
 - race conditions cause bugs that are hard to reproduce!
- Besides race conditions, another issue is **synchronization** (one process is waiting for the results computed by another)
 - can we avoid busy waiting?

Say you want to count the total number of squirrels and birds...

- Matt: squirrel enthusiast
 - He gets excited (and increments a counter) whenever he sees a squirrel.
 - He ignores birds.
- Jonathan: bird enthusiast
 - He gets excited (and increments a counter) whenever she sees a bird.
 - He ignores squirrels



Matt & Jon have threads that share the same memory

- Two threads in loop incrementing a counter
 - Matt increments *event* when a squirrel arrives

```
while (TRUE) {  
    wait_for_squirrel();  
    event = event + 1;  
}
```
 - Jonathan increments *event* when a bird arrives

```
while (TRUE) {  
    wait_for_bird();  
    event = event + 1;  
}
```

Compiling “event=event+1”

Squirrel-watcher thread:

1. move event, R0
 2. increment R0
 3. move R0, event
- R0 is used as a temporary register for “event”

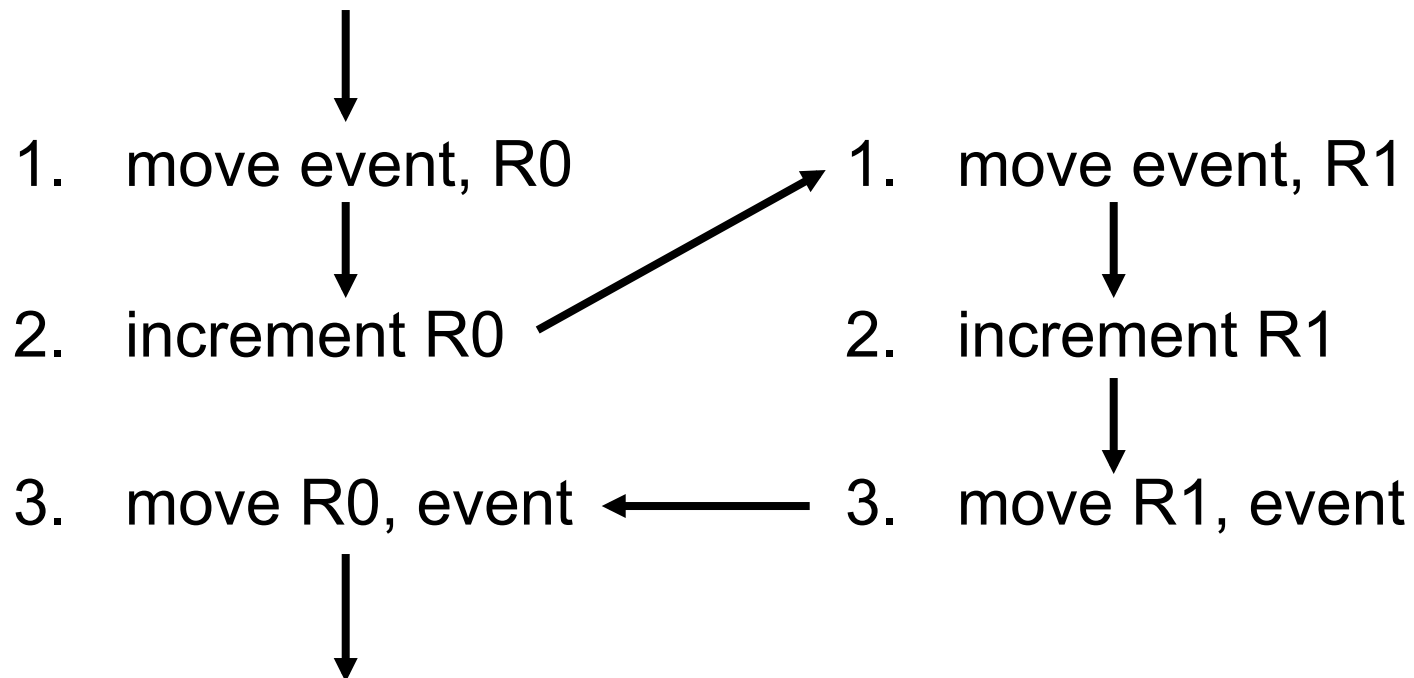
Bird-watcher thread:

1. move event, R1
 2. increment R1
 3. move R1, event
- notice how we were careful to avoid using the same register
 - But does it work?

Does this work if squirrels and birds arrive at the same time?

Squirrel-watcher thread:

Bird-watcher thread:



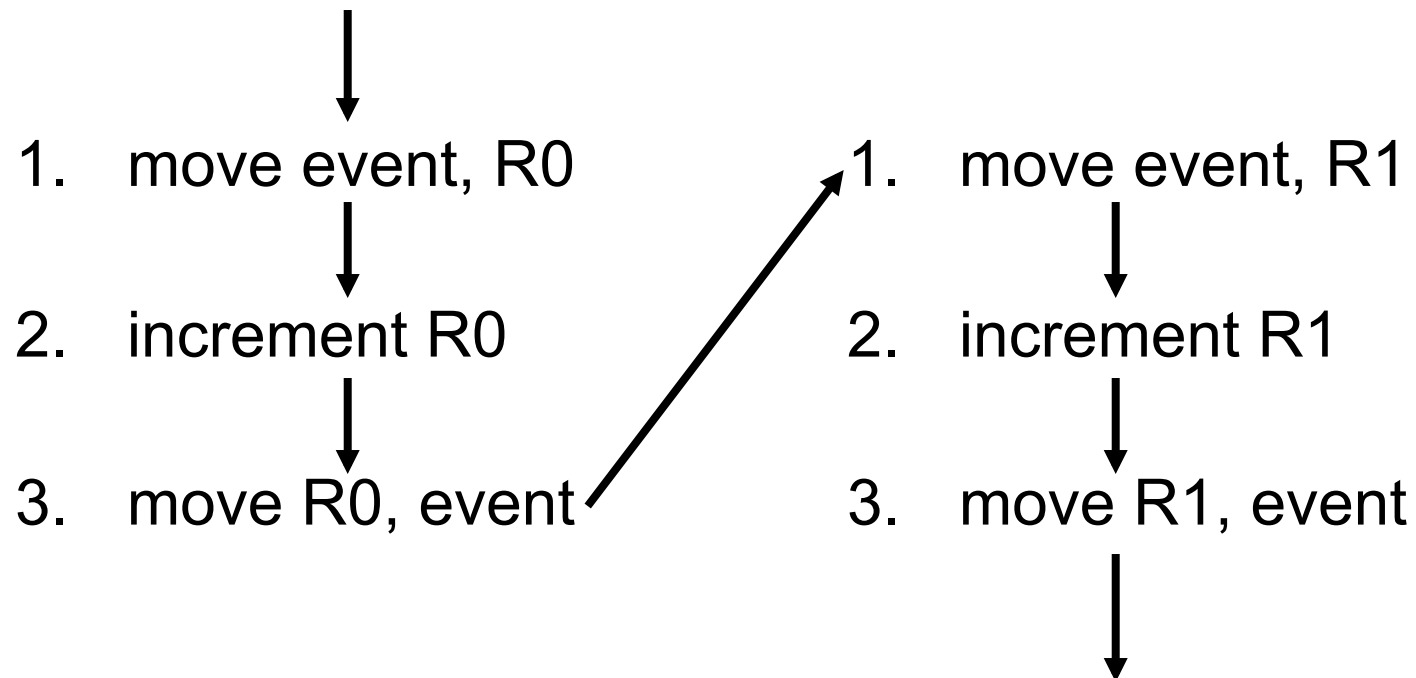
Reasonable questions to ask at this point...

- At the end of this interleaving, by how much does *event* get incremented?
- Are there other “incorrect” interleavings of these instructions?
- Are there *any* “correct” interleavings?
 - if so, how can we ensure that they occur?

An alternative interleaving:

Squirrel-watcher thread:

Bird-watcher thread:



What's going on here?

- Things work as expected if these three instructions are executed *together*:
 1. move event, R1
 2. increment R1
 3. move R1, event
- Other threads that operate on *event* shouldn't interrupt during this operation
- We'd like a way to “group” these three instructions together, so that other relevant threads/processes won't interfere with them
- This is called the *mutual exclusion problem*

In other words...

1. BEGIN CRITICAL SECTION
 - (no one should interrupt)
2. move event, R1
3. increment R1
4. move R1, event
5. END CRITICAL SECTION
 - (OK, interrupt again)


Mutual exclusion

- Need to ensure “atomic” execution of a sequence of instructions
 - at least as far as all the other threads accessing the data are concerned
- How might we do this?
 - an interrupt might occur at any time, giving control to another thread

Possible ways to implement

- Turn off interrupts
 - no help from OS or CPU
 - kind of drastic
- “TEST AND SET” instruction
 - “spin lock”
 - help from CPU
 - can be expensive
- Block until you can have exclusive access
 - Some kind of system call
 - help from OS

Classic example: Producer/Consumer Problems



- from time to time, the producer places an item in the buffer
- the consumer removes an item from the buffer
- careful synchronization required (they run simultaneously)
- the consumer must wait if the buffer empty
- the producer must wait if the buffer full
- typical solution would involve a shared variable **count**
- also known as the Bounded Buffer problem
- Example: in UNIX shell

cat myfile.txt | lpr

Hierarchy of Abstractions

Idealized Problems
Producer-Consumer
Dining Philosophers
Readers-Writers

High-level Synchronization Primitives
Monitors (Hoare, Brinch-Hansen)
Synchronized method in Java

OS-level support (mutual exclusion and synchronization)
Special variables: Semaphores, Mutexes
Message passing primitives (send and receive)

Low-level (for mutual exclusion)
Interrupt disabling
Using read/write instructions
Using powerful instructions (Test-and-set, Compare-and Swap...)

What on earth did that mean?

- Computer scientists like to find ways to think about problems at different *layers*
 - this is actually useful sometimes (not just as employment program for CS PhDs)
- Toward the bottom there are low-level *mechanisms*
 - use these to build a solution
- At the top there are high-level *problems*
 - you can try to adapt a solution to an *existing* standard problem to work for *your* problem
 - just find a way to “reduce” your problem to the standard problem

Dealing with Synchronization & Mutual Exclusion

- From the top: find similarities between problem at hand and high-level “ideal” problems with “standard” solutions
 - dining philosophers, producer/consumer, etc
- From the bottom: toolkit of low-level “primitives” that address various aspects
 - test and set instructions, disable interrupts, etc.
- In the middle: *abstractions* that link the two
 - monitors, semaphores, message-passing
 - OS and language interfaces to lower-level tools

The *Mutual Exclusion* Problem

- Problem: Allow access to shared data structures without *race conditions* caused by instruction interleaving
- Abstract Solution:
 - Identify *critical sections* (aka *critical regions*)
 - e.g., instruction sequence for “event = event + 1”
 - Give critical sections exclusive access (without “interference”) over their entire execution
- Several different “standard” ways to implement critical section abstraction

Critical Section Abstraction

- Like a “guard” protecting regions of code:
 - BEGIN CRITICAL SECTION
 - if no one else in critical section, you can go in
 - otherwise wait
 - do critical section stuff
 - END CRITICAL SECTION
 - relinquish exclusive access, let someone else enter
- Assumption: *everyone else* is also playing by the same rules – all processes include explicit code around their critical sections

Requirements for a good critical section solution

- Safety
 - should really work - no two threads should ever simultaneously be in a critical section
- Generality
 - shouldn't depend on “fragile” assumptions
- Deadlock Freedom
 - no state where *everyone* is waiting for someone *else* to do something before *anyone* can proceed
- Starvation Freedom (“bounded liveness”)
 - if a thread is waiting for access to a critical section, it should eventually be granted

Approach #1:

Turning off interrupts

- Works by preventing pre-emptive scheduling via a clock or other interrupt
 - turn interrupts off at beginning, on again at end
- Used in low-level parts of the OS (e.g., during interrupt handling)
- Meets requirements, but not great for user processes
 - requires care: failure to yield may require reboot
 - overly powerful: prevents *all* other process from preempting, not just those entering critical sections
 - inflexible: all-or-nothing

Approach #2a:

Use shared variables

- GL Peterson's solution (1982)
- Everyone needing access to critical section shares special variables used to “flag” access
- Assumes that simple assignments to, and tests on, the shared variables are *atomic*
 - e.g., V doesn't change *during* “ $V=1$ ”
 - V might change right before or after, however
 - this is a pretty safe assumption on most uniprocessors

Peterson's solution

- Three shared Boolean variables
 - turn, flag[2] --- (supports 2 threads)
- Code to **enter** critical section:
flag[ME] = TRUE
turn = ME
while (turn==ME) && (flag[1-ME]==TRUE)
; /* busy wait */
- Code to **exit** critical section:
flag[ME] = FALSE

Is it really that hard?
What if we just do this:

```
while (turn != ME)
    ; /* busy waiting */
CriticalSectionHere();
turn = 1-ME; /* be fair to other */
```

Ensures mutual exclusion, but requires *strict alternation*.
A process can't ever enter its CS twice in succession
if the other process doesn't enter CS!

How about this?

```
while (flag[1-ME])  
    ; /* wait if other guy in CS */  
flag[ME] = TRUE; /* declare your entry */  
CriticalSectionHere();  
flag[ME] = FALSE; /* unblock other guy */  
Non_CS();
```

Safety requirement violated (race condition)!

- P0 tests flag[1] and finds it False

- P1 tests flag[0] and finds it False

- Both proceed, set their flags to True, and enter CS

Or this???

```
flag[ME] = TRUE; /* declare entry first */
while (flag[1-ME])
    ; /* wait while other is guy in CS */
CriticalSectionHere();
flag[ME] = FALSE; /* release */
Non_CS();
```

Vulnerable to deadlock (not deadlock-free)!

P0 sets flag[0] to TRUE

P1 sets flag[1] to TRUE

Both enter their loops and keep waiting

Sorry!

All that complexity in
Peterson's solution
is actually necessary

Peterson's Solution (starting to look good...)

```
flag[ME] = TRUE; /* declare interest */
turn = ME; /* for race condition */
while ((flag[1-ME]==TRUE)
        && (turn == ME))
    ; /* busy wait */
CriticalSectionHere();
flag[ME] = FALSE; /* release */
Non_CS();
```

other guy(1-ME)
is contending

detect and deal with
race condition

Approach 2b:

Use the hardware (TSL)

- Peterson's solution is nice, but complex
- Much simpler when we have CPU support
 - reading and writing in one atomic operation
 - and fortunately, most modern CPUs have something like this
- Test-and-set-lock: TSL R, memory
 - R gets content of memory
 - memory gets value “1”
- Other powerful atomic instructions also work
 - swap, compare-and-swap, load-linked

Mutual Exclusion with the TSL instruction

- Enter critical section:
 1. TSL R0, lock
 2. if R0 == 1 jump to “1”
- Leave critical section:
 1. lock = 0;
- Does this work? Why (or not)?

So far:

- Solution 1: Disable interrupts
 - big hammer
- Solution 2: busy wait with shared variables
 - 2a: Peterson's solution
 - complex but no hardware support
 - works on virtually all uniprocessors
 - 2b: TSL
 - simpler but assumes TSL (or equiv) instruction
 - works even on multiprocessors
 - Can be generalized to > 2 threads

Are we done yet?


- Not quite!
- The shared variable solutions require *busy waiting*
 - bad: this can be inefficient - the locked out thread still needs to consume CPU while waiting
 - worse: can deadlock if the waiting process has higher priority in the scheduler
- We'd usually prefer a *blocking* solution

Approach 3:

Sleep and Wakeup

- Two *abstract* system calls (these are not the names of actual Unix system calls!):
 - **sleep**
 - blocks until someone calls **wakeup**
 - **wakeup**(*process*)
 - unblocks *process*
 - process should have been previously blocked with **sleep**
- Remember the producer-consumer problem?
 - sleep and wakeup can help solve the full / empty buffer problem

Producer/Consumer Problems



- from time to time, the producer places an item in the buffer
- the consumer removes an item from the buffer
- careful synchronization required (they run simultaneously)
- the consumer must wait if the buffer empty
- the producer must wait if the buffer full
- typical solution would involve a shared variable count
- also known as the Bounded Buffer problem
- Example: in UNIX shell

cat myfile.txt | lpr

Producer/Consumer (partial solution)

- Producer:

```
while (TRUE) {  
    produce X  
    if (count == N)  
        sleep;  
    add_to_buffer(X)  
    count = count + 1;  
    if (count == 1)  
        wakeup(Consumer);  
}
```

- Consumer:

```
while (TRUE) {  
    if (count==0)  
        sleep;  
    remove_from_buffer(X)  
    count = count -1;  
    if (count == N-1)  
        wakeup(Producer);  
    consume X  
}
```

Doesn't quite work

- Count is initially 0
 - Consumer reads the count
- Producer produces the item, inserts it, and increments count (to 1)
- Producer executes **wakeup**, but there is no waiting consumer (at this point)
- Now consumer continues its execution and calls **sleep**; consumer blocks
- Consumer stays blocked forever
 - Main problem: race condition -- *wakeup* was lost

Semaphores (Dijkstra)

- A semaphore **S** has a non-negative integer value
- Two operations
 - **up(S)** (aka V(S)) : increments the value of S
 - **down(S)** (aka P(S)) : decrements the value of S if S is positive, else makes the calling thread/process wait
- When $S=0$, down(S) moves the thread to sleep (blocked) state
 - no busy waiting
- If $S=0$, up(S) also wakes up one sleeping process (if there are any)
 - uses an internal list of sleeping processes
- up and down calls are *atomic* actions

Can we do mutual exclusion via Semaphores?

- Semaphore value S initially set to 1
- Enter critical section
 - $\text{down}(S)$
- Exit critical section
 - $\text{up}(S)$
- Does this work?

Neat tricks with Semaphores

- Simplest semaphore examples are binary (semaphore value S is either 1 or 0),
 - used here as an “improved” (blocking) TSL
 - to get strict mutual exclusion, set initial $S=1$
 - maximum number in critical section = 1
- Initial values of $S > 1$ can be used to allow an arbitrary maximum number of threads to be active somewhere
 - other applications besides mutual exclusion (e.g., resource load control)

Mutual Exclusion Toolkit

- Solution 1: Disable interrupts
 - usually too big a hammer
- Solution 2: busy wait with shared variables
 - exploits standard instruction set features
 - 2a: Peterson's solution
 - complex but no hardware support needed
 - works on virtually all uniprocessors
 - 2b: TSL
 - simpler but assumes TSL (or equiv) instruction
 - works even on multiprocessors
- Solution 3: blocking system calls
 - needs OS support
 - 3a sleep/wakeup
 - simple; useful building block
 - “lost wakeup problem”
 - 3b Semaphores
 - generalized abstraction
 - can solve lost wakeup problem
 - can be built atop interrupt disabling and busy waiting inside OS