

# CIS 505 – Software Systems

Notes for 1/26/2012

Matt Blaze

# Our problem:

- Resources (such as memory) shared by  $> 1$  thread
  - things can get amazingly (and subtly) screwed up because of, among other things, *race conditions*
    - strange answers, deadlocks, etc
- So we need some way to “keep order”
  - mutual exclusion
  - thread synchronization
- Maybe we can solve in the OS, or maybe in the applications, or maybe both

# So far:

- Solution 1: Disable interrupts
  - big hammer
- Solution 2: busy wait with shared variables
  - 2a: Peterson's solution
    - tricky (3 variables), but requires no hardware support
    - works on virtually any conventional processor
  - 2b: TSL
    - simpler, but assumes TSL (or equiv) instruction
    - works even on multiprocessors
  - Can be generalized to  $> 2$  threads

# Limitations of Peterson/TSL


- Busy waiting!
  - “are we there yet?” in a tight loop
  - that is, the waiting thread doesn't *block*
- Designed specifically for mutual exclusion
  - you need other stuff to do general synchronization on top of them
- But, at least, they're very general
  - no help from OS, so useful *inside* OS

# Approach 3:

## Sleep and Wakeup

- Thread blocking and unblocking synchronization
- Two abstract system calls (*not* actually what they're called in Unix systems!):
  - sleep()
    - blocks until someone calls wakeup
  - wakeup(process)
    - unblocks process, which should be blocked in sleep
- Remember the producer-consumer problem?
  - this helps solve the full / empty buffer problem
  - problem here is *synchronization*, not just mutual exclusion

# Producer/Consumer Problems



- from time to time, the producer places an item in the buffer
- the consumer removes an item from the buffer
- careful synchronization required (they run simultaneously)
- the consumer must wait if the buffer empty
- the producer must wait if the buffer full
- typical solution would involve a shared variable count
- also known as the Bounded Buffer problem
- Example: in UNIX shell

**cat myfile.txt | lpr**

# Producer/Consumer (partial solution)

- Producer:

```
while (TRUE) {  
    produce X...  
    if (count == N)  
        sleep;  
    add_to_buffer(X)  
    count = count + 1;  
    /* need above to be atomic */  
    if (count == 1)  
        wakeup(Consumer);  
}
```

- Consumer:

```
while (TRUE) {  
    if (count==0)  
        sleep;  
    remove_from_buffer(X)  
    count = count -1;  
    /* need above to be atomic... */  
    if (count == N-1)  
        wakeup(Producer);  
    consume X...  
}
```

# Doesn't quite work

- Count is initially 0
  - Consumer reads the count
- Producer produces the item, inserts it, and increments count (to 1)
- Producer executes wakeup, but there is no waiting consumer (at this point)
- Now consumer continues its execution and goes to sleep
- Consumer stays blocked forever
  - Main problem: wakeup was lost



# A new mechanism: Semaphores (Dijkstra)

- A semaphore **S** has a non-negative integer value
- Two operations
  - **up(S)** (aka V(S)) : increments the value of S
  - **down(S)** (aka P(S)) : decrements the value of S if S is positive, else makes the calling thread/process wait
- When  $S=0$ , down(S) moves the thread to sleep (blocked) state
  - no busy waiting
- If  $S=0$ , up(S) also wakes up one sleeping process (if there are any)
  - uses an internal list of sleeping processes
- up and down calls must be *atomic* actions

# Can we do mutual exclusion with Semaphores?

- Yep, we sure can:
  - Set semaphore S with initial value of 1
  - To enter critical section
    - `down(S)`
  - To exit critical section
    - `up(S)`
- That's it.
  - Does this really work? Why?

# Neat tricks with Semaphores

- Simplest semaphore examples are binary (semaphore value  $S$  is either 1 or 0),
  - used here as an “improved” (blocking) TSL
  - to get strict mutual exclusion, set initial  $S=1$ 
    - maximum number in critical section = 1
- Initial values of  $S > 1$  can be used to allow an arbitrary maximum number of threads to be active somewhere
  - other applications besides mutual exclusion (e.g., resource load control)

# Basic Semaphore Implementation

```
typedef struct {  
    int value;  
    *pid_t wait_list; /* list of processes */  
} semaphore;
```

```
down(semaphore S){  
    if (S.value >0) S.value--;  
    else { add this process to S.wait_list;  
          sleep;  
    }
```

```
up(semaphore S){  
    if (S.wait_list==null) S.value++;  
    else { select and remove a process P from S.wait_list;  
          wakeup(P) ;  
    }
```

**up** and **down** have to be protected as critical sections

# Producer-Consumer, with semaphores

```
semaphore:  full = 0      /* number of full slots */  
            empty = SLOTS /* number of empty slots */  
            mutex = 1     /* critical section (binary) */
```

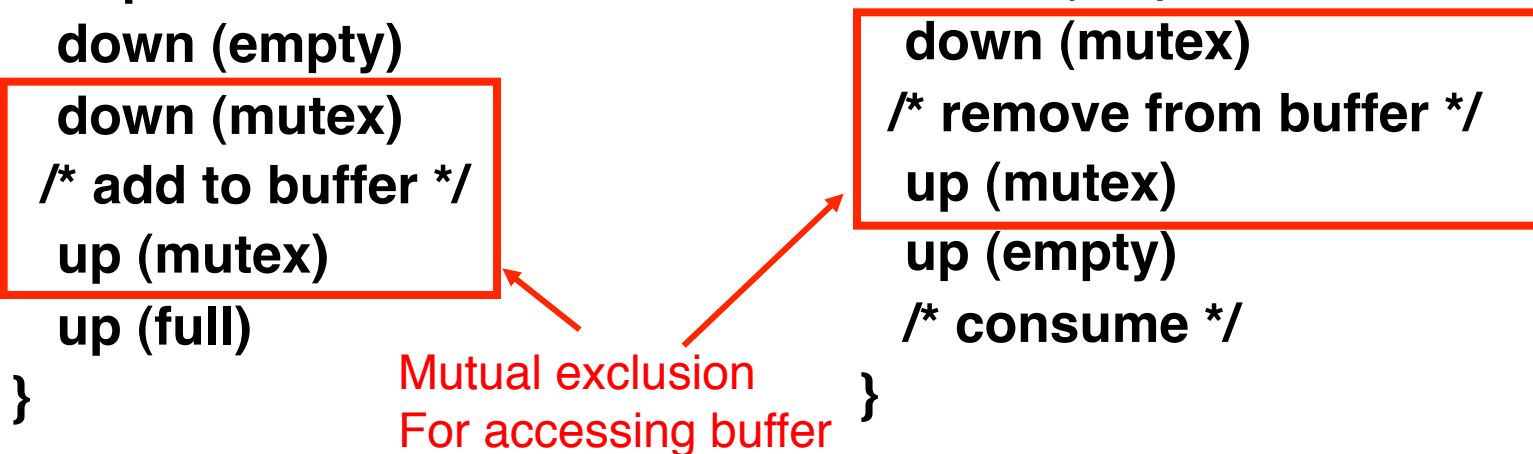
**Producer code:**

```
while (TRUE) {  
    /* produce */  
    down (empty)  
    down (mutex)  
    /* add to buffer */  
    up (mutex)  
    up (full)  
}
```

**Consumer code:**

```
while (TRUE) {  
    down (full)  
    down (mutex)  
    /* remove from buffer */  
    up (mutex)  
    up (empty)  
    /* consume */  
}
```

Mutual exclusion  
For accessing buffer



# Producer-Consumer

## does it work yet?

```
semaphore:  full = 0      /* number of full slots */  
            empty = SLOTS /* number of empty slots */  
            mutex = 1     /* critical section (binary) */
```

**Producer code:**

```
while (TRUE) {  
    /* produce */  
    down (empty)  
    down (mutex)  
    /* add to buffer */  
    up (mutex)  
    up (full)  
}
```

**Consumer code:**

```
while (TRUE) {  
    down (full)  
    down (mutex)  
    /* remove from buffer */  
    up (mutex)  
    up (empty)  
    /* consume */  
}
```

What happens if we switch the order of down(empty) and down(mutex) ?  
What happens if we switch the order of up(mutex) and up(full) ?

Making it work in practice:  
semaphores in real systems

# Semaphores in actual UNIX/POSIX systems

- There aren't system calls called *up* or *down*
  - naturally, it's more complicated than that
  - semaphores have traditionally been an internal feature of the UNIX kernel, but not available to user processes
- Instead, there's a family of POSIX standard system calls that implement semaphores
  - one data type: `sem_t` data
  - five system calls: `sem_init`, `sem_destroy`, `sem_wait`, `sem_trywait`, `sem_post`




# POSIX Semaphore Syscalls

- *sem\_t*: data type for semaphores
- *int sem\_init(sem\_t \*sp, unsigned int count, int type)*: Initialize semaphore pointed to by sp to count. (*type* controls details of semaphore behavior)
- *int sem\_destroy(sem\_t \*sp)*: destroys any state related to the semaphore pointed to by sp.
- *int sem\_wait(sem\_t \*sp)*: blocks the calling thread until the semaphore count pointed to by sp is greater than zero, and then it atomically decrements the count.
- *int sem\_trywait(sem\_t \*sp)*: atomically decrements the semaphore count pointed to by sp if it is greater than zero; otherwise, it returns an error.
- *int sem\_post(sem\_t \*sp)*: atomically increments the semaphore count pointed to by sp. If there are any threads blocked on the semaphore, one is unblocked.

# The infamous Dining Philosophers...

- Assumptions:
  - philosophers can only afford to eat rice
  - worse, they only have one chopstick each!
    - but need two to eat
  - alternate between two states:
    - complaining (or looking for work)
    - trying to eat
- Sadistic computer scientist arranges them in circle:



# Dining Philosophers?

- Philosophers have good table manners
  - must acquire two chopsticks to eat
- Useful abstraction / “standard problem”:
  - soundness
    - every chopstick is held by  $\leq 1$  philosopher at a time
  - deadlock-freedom
    - no state where no one gets to eat
  - starvation-freedom
    - solution guarantees that *all* philosophers occasionally eat

# First try at a solution

- Obvious solution is to represent each chopstick by a semaphore (semaphore chopstick[N]).
- Down before picking it up & up after using it.

```
while (TRUE) { /* philosopher i loop */
    DOWN(chopstick[i]);
    DOWN(chopstick[i+1 mod N]);
    eat();
    UP(chopstick[i]);
    UP(chopstick[i+1 mod N]);
    lookforjob();
    complain();
}
```

# Did that work?

- No

# What was wrong?

- Solution is *sound*
  - the semaphores indeed guarantee that every chopstick is in use by  $\leq 1$  philosopher
- But vulnerable to *deadlock*
  - Bad scenario: everyone runs in lockstep
  - first everyone picks up their chopstick[i]
    - that works
  - then everyone tries to get chopstick[i+1 mod N]
    - but that's someone else's chopstick[i], which is in use
    - so they all block
  - kaboom. deadlock.

# Possible solutions (do they work?)

- ? Make get chopstick[i] and get chopstick[i+1 mod N] atomic by surrounding with a binary (“mutex”) semaphore ?
- ? Reverse the order of get chopstick[i] and get chopstick[i+1 mod N] for every other philosopher ?
- ? Forget the chopstick semaphores ?
  - ? protect eat() with mutex semaphore ?

# A better solution

- Add a binary mutex semaphore
- Each philosopher in one of three states
  - *hungry* (wants to eat), *eating* (holding 2 chopsticks) or *complaining*
- To eat, add this code:
  - down(mutex)
  - state[i]=hungry
  - state[i]=eating if neither neighbor is eating
  - up(mutex)
- To finish eating, if either neighbor is hungry, up() their chopstick
- See Tannenbaum figure 2-33 for details



# Some “standard” problems (learn them!)

- See Tanenbaum, *Modern Operating Systems*, ch 2
- Dining philosophers
- Readers and Writers
  - access to a shared database
  - deadlocks, starvation, consistency
- Sleeping Barber
  - one barber, one chair, n chairs in waiting area
    - barber sleeps when no customers
    - customers arrive asynchronously
  - queuing, race conditions, starvation, deadlock

# Mutual Exclusion Toolkit

- Solution 1: Disable interrupts
  - big hammer
- Solution 2: busy wait with shared variables
  - 2a: Peterson's solution
    - complex but no hardware support required
    - works on virtually all uniprocessors
  - 2b: TSL
    - simpler but assumes TSL (or equiv) instruction
    - works even on multiprocessors
- Solution 3: blocking system calls
  - 3a sleep/wakeup
    - simple; useful building block
    - “lost wakeup problem”
  - 3b Semaphores
    - generalized abstraction
    - can solve lost wakeup problem
  - can be built atop interrupt disabling and busy waiting inside OS

OK, new topic:

Deadlocks

# Deadlocks and Resource Allocation

- The deadlock-vulnerable environment
  - finite *resources*, each of which can be assigned to a limited number of processes/threads (usually at most one)
  - processes can request multiple resources
    - several at once or one at a time
    - allowed to accumulate more while holding
    - block until requested resource is available
  - processes can release resources when finished with them
- A *deadlock* is a situation where there is a set of blocked processes waiting for resources such that there is no way to satisfy any of their requests
  - even if all the *unblocked* processes release theirs
- We've already encountered deadlocks
  - dining philosophers, etc

# Our (not-so-great) deadlock toolkit so far

- Approach 1: Ignore problem
- Approach 2: Prevention
  - 2a: Exhaustive Search of Possible States
  - 2b: Ad Hoc Genius
- Approach 3: Detection
  - 3a: Ad Hoc Genius
- We'll be focusing on Prevention, Detection
  - making our toolkit more systematic
- Also a new idea: *Deadlock Avoidance*

# Approach 1: Ignore Problem

- Deadlocks aren't always so bad
  - sometimes we can just live with the possibility
  - when things seem to stop, just start over
  - fine for many applications
- But unfortunately, sometimes they *are* bad
  - inside the OS, databases, transaction systems
- Usual name: “Ostrich Algorithm”
  - “stick head in sand”
  - (but real ostriches aren't actually that dumb)

# Approach 2:

## Deadlock Prevention

- 2a: Ad hoc approach
  - uses ingenious insight to notice potential deadlocks
  - works well only for very smart people
  - example: dining philosophers
- 2b: Exhaustive search approach
  - go through every possible state, look for deadlocks
  - exponential, gets too much to deal with quickly
- We'd prefer something more systematic

# Approach 3:

## Deadlock Detection & Recovery

- Ad hoc approach
  - notice that a deadlock has occurred, then wildly panic
    - start killing off processes until things start to move again
  - have to be careful that killed off processes can be rolled back / restarted safely
  - may have to reboot OS
- We'd prefer something more systematic



# Approach 2, redux: Deadlock Prevention

(a more systematic approach)

# The four preconditions for deadlock

- #1 *Mutual Exclusion*
  - resources being held by one process can't simultaneously be used by another
- #2 *No resource preemption*
  - holders keep resources until relinquished
- #3 *Hold and wait*
  - resource holders can “accumulate” (request additional resources and block waiting for them)
- #4 *Circular waiting*
  - circular chain of two (or more) processes/threads
  - each waiting for resource held by next in chain

# The *entire secret* to deadlock prevention:

- Just eliminate *any one* of the four preconditions
  - doesn't matter which one!
  - if you do this, deadlock can't possibly happen, ever!
  - problem solved!
- That's all there is to it!
  - but... how might we do this?

# Removing Precondition 1:

## *Mutual Exclusion*

- Mutual Exclusion: Resources held by one process can't simultaneously be used by another
- This is hard to avoid -- usually you either need mutual exclusion or you don't
- Sometimes mutex can be avoided by making resource use effectively "atomic"
  - use resource in a burst of exclusive access
  - spooling

# Removing Precondition 2:

## *No Resource Preemption*

- No Resource Preemption: Resource holders keep their resources until they relinquish them
- Maybe can eliminate with some kind of policy
  - e.g., if high priority process wants a resource, it gets it and the resource is taken away from lower priority
- But no good way to implement this in practice
  - complex recovery when resource is taken away

# Removing Precondition 3:

## *Hold and Wait*

- Hold and Wait: Resources can be accumulated (requested and blocked for while holding others)
- Maybe can eliminate by requiring complete set of resources to be requested together
  - must release all to request new set
- Inefficient
  - and what about resources whose identity can't be determined in advance (e.g., file names)?

# Eliminating Hold-and-Wait: Two-Phase Locking

- Processes manage resources in two phases
  - locking
    - attempt to acquire locks
    - if any attempt would block, release and start over
  - using
    - use any resource only after all locks acquired
    - release locks when finished
- Safe, but still inefficient for many applications
  - similar to spin locks
  - used occasionally in database

# Removing Precondition 4:

## *Circular Waiting*

- Circular Waiting: circular chain of processes, each waiting for resource held by next process in chain
- May be able arrange requests to break cycles
  - this better work, it's the only prevention idea left
- Central idea: **Hierarchical Allocation**
  - number resources; only allow obtaining resources with higher numbers than those currently held
  - requires finding a suitable numbering



# Hierarchical Allocation and Dining Philosophers

- Obvious solution is to represent each chopstick by a semaphore (semaphore chopstick[N]).
- Down before picking it up & up after using it.

```
while (TRUE) { /* philosopher i loop */
    DOWN(chopstick[i]);
    DOWN(chopstick[i+1 mod N]);
    eat();
    UP(chopstick[i]);
    UP(chopstick[i+1 mod N]);
    lookforwork()
    complain();
}
```

## Now, what about this...

- ? Reverse the order of get chopstick[i] and get chopstick[i+1 mod N] for every other philosopher ?
  - what does this do?

or....

- if  $i+1 \bmod N == 0$  , get chopsticks in reverse order
  - fixes deadlock (not starvation, though)

# Bottom line on Deadlock Prevention

- It sounds so easy!
  - you only have to eliminate any *one* of the four deadlock preconditions
- But eliminating even one of the preconditions is often hard
  - *Hold-and-Wait* (#3) can sometimes be eliminated with two-phase locking
  - *Circular waiting* (#4) can sometimes be eliminated in practice through Hierarchical Allocation

# Our improved deadlock toolkit (so far)

- Approach 1: Ignore problem
- Approach 2: Prevention
  - 2a: Exhaustive Search of Possible States
  - 2b: Ad Hoc Genius
  - **2c: New! Two-Phase Locking**
  - **2d: New! Hierarchical Allocation**
- Approach 3: Detection and Recovery
  - 3a: Ad Hoc Genius
  - **3b: ??? coming up next ???**
- Approach 4, also still to come: Deadlock *Avoidance*

# Approach 3 Revisited: Deadlock Detection/Recovery

(a more systematic approach)

# Deadlock Detection

- Even if you can't prevent deadlocks, perhaps you can systematically *detect* them
  - should be coupled with a systematic recovery scheme
- So what we need are
  - algorithms for detecting deadlocks
    - graph theory to the rescue
  - a way of “undoing” what we did before the deadlock occurred
    - “rollback”

# Are there good deadlock detection algorithms?

- Basic idea: represent resource allocation as a graph, look for cycles
  - no cycles == no deadlocks
  - cycles == deadlocks
- Cycle detection! Sounds like a typical graph theory problem
  - can this be done efficiently?
- Also a systems problem
  - when to check?

# Detecting Deadlocks (simple case)

- Example 1: Is this a deadlock?
  - P1 has R2 and R3, and is requesting R1
  - P2 has R4 and is requesting R3
  - P3 has R1 and is requesting R4
- Example 2: Is this a deadlock?
  - P1 has R2, and is requesting R1 and R3
  - P2 has R4 and is requesting R3
  - P3 has R1 and is requesting R4
- Solution: Build a graph: Resource Allocation Graph (RAG)
  - There is a node for every process and a node for every resource
  - If process P currently has resource R, then put an edge from R to P
  - If process P is requesting resource R, then put an edge from P to R
- There is a deadlock if and only if RAG has a cycle



# Resource Allocation Graph Cycle Detection Algorithms

- Can be generalized to multiple instances of resources
- A little slow (quadratic) but still often feasible
- Still leaves us with the problem of *recovery...*

# After Detection: Deadlock Recovery

- Preemption
  - Take away a resource from current owner
  - Frequently impossible
- Rollback
  - Checkpointing periodically to save states
  - Reset to earlier state before acquiring resource
  - Used in database systems
- “Soprano’s” Algorithm: Manual Killing
  - Crude but simple
  - Keep killing processes in cycle until cycle is broken

# Our deadlock toolkit (so far)

- Approach 1: Ignore problem
- Approach 2: Prevention
  - 2a: Exhaustive Search of Possible States
  - 2b: Ad Hoc Genius
  - **2c: New! Two-Phase Locking**
  - **2d: New! Hierarchical Allocation**
- Approach 3: Detection and Recovery
  - 3a: Ad Hoc Genius
  - **3b: New! Cycle Detection (with recovery)**
- Now: Approach 4: Deadlock *Avoidance*

# Approach 4: Deadlock Avoidance

(another approach)

# *Preventing* Deadlock vs. *Avoiding* Deadlock

- Deadlock *prevention* depends on designing the system/application to deny at least one of the preconditions for deadlock
  - advantage: won't ever deadlock no matter what, e.g., the OS scheduler does
  - disadvantage: may not be possible/efficient for some kinds of resources
- Another approach, called deadlock *avoidance*, asks the OS to schedule resource use (at *runtime*) in a “safe” way
  - “unsafe” schedules are those that risk deadlock

# Safe States and Deadlock

- Environment: there are *multiple, interchangeable* instances of resources
  - example: memory
  - this is *not* the case with, e.g. mutexes, chopsticks in DP
  - processes can state requirements in advance
  - system free to make choices at runtime about which resources to allocate (and who to unblock when)
    - but no preemption - once resource given, it can't be taken back until the process gives it back
- “Safe” states result from making these allocation choices *conservatively*
  - never allocate in a way that could result in deadlock

# (Very) Conservative Allocation (fundamentalist, in fact)

Example: total resources: 12

process	maximum requirement
A	6
B	11
C	7

Looks like we can only run one at a time, no matter how many resources are actually used by each process at any given moment.

Too bad! But maybe that's just the price you pay for being so risk averse...

# How conservative is *too* conservative here?

- It may be possible to be conservative in allocating resources without assuming that everyone is always using their maximum
  - is there a more *efficient* way to be conservative?
- Consider a bank that issues lines of credit:
  - greedy: wants as many customers as possible
    - customers will eventually repay their debts, but may need to max out their credit line first
      - customers willing to wait to get their loans
    - total lines of credit may exceed bank's actual assets (but actual money loaned out is always  $\leq$  total assets
  - conservative: must never, ever get deadlocked



# Banker's (Dijkstra's) Algorithm

- Environment:  $n$  process  $P_1, \dots, P_n$  and  $m$  resources  $R_1 \dots R_m$
- Every process declares (in advance) its *claim*---the maximum number of resources it will ever need
  - Sum of claims of all processes could exceed total number of resources
- To avoid deadlocks, OS maintains the *allocation state*
  - Current allocation matrix  $C$ :  $C[i,j]$  is the number of instances of resource  $R_j$  currently held by process  $P_i$
  - Claims matrix  $M$ :  $M[i,j]$  is the maximum number of instances of  $R_j$  that process  $P_i$  will ever request
  - Availability vector  $A$ :  $A[j]$  is the number of instances of  $R_j$  currently free.
- Suppose process  $P_i$  requests certain number resources. Let  $Req$  be the request vector ( $Req[j]$  is number of requested instances of  $R_j$ )
  - Valid request if  $Req \leq M[i]-C[i]$  (i.e. it should be in accordance with claim)
  - If  $Req \leq A$ , then it is possible for OS to grant the request
  - Avoidance strategy: Deny the request if the resulting state will be *unsafe*

# Safe states

- An allocation state is **safe** if there is an ordering of processes (a **safe sequence**) such that:
  - the first process can finish for sure
    - there are enough available resources to satisfy all of its claim
  - once the first process releases its resources, the second process can finish for sure (even if it asks all its claim)
  - and so on.
- This is “safe” because the OS always can avoid deadlock simply by blocking *new* requests until some or all of the current processes have finished
  - check to see if a new request would be unsafe
    - if so, block it until someone finishes, then check again

# Example

(One resource class only)

total resources: 12

unallocated: 2

process	holding	max claims
A	4	6
B	4	11
C	2	7

safe sequence: A,C,B

# Another example

(One resource class only)

total resources: 12

unallocated: 2

process	holding	max claims
A	4	6
B	4	11
C	2	9

safe sequence: none!

# Banker's algorithm

- Maintain claims  $M$ , current allocation  $C$  and current availability  $A$
- Suppose process  $P_i$  requests  $Req$  such that  $Req \leq A$  and  $Req + C[i] \leq M[i]$ 
  - Consider the state resulting from granting this request (i.e. by adding  $Req$  to  $C[i]$  and subtracting  $Req$  from  $A$ ).
  - Check if the new state is a safe state. If so, grant the request, else deny it.

# Checking Safety

- How do we check if an allocation state is safe?
  - Current allocation matrix  $C$
  - Maximum claims matrix  $M$
  - Availability vector  $A$
- Same as running a deadlock detection algorithm assuming that every process has requested maximum possible resources
  - Choose Requests Matrix  $R$  to be  $M - C$ , and see if the state is deadlocked (is there an order in which all of these requests can be satisfied).

# Variations on the Banker's Algorithm

- Multiple resource types
  - just run it for each resource
- “increased” and “decreased” claims at runtime

# When is Deadlock Avoidance practical?

- Requires that processes know (and state) their maximum requirements *up front*
  - even if not using them all at once
  - like a “credit line”
- But general processes may not be able to do that for most of the kinds of resources they use
- Useful mainly for
  - specialized applications that share a few classes of inter-changeable resources
  - OS-based resources like memory & processor allocation



# Evaluating our deadlock toolkit

- Approach 1: Ignore
- Approach 2: Prevention
  - System design rules
  - 2a: Exhaustive Search of Possible States
  - 2b: Ad Hoc Genius
  - **2c: New! Two-Phase Locking**
    - inefficient (spin locks)
  - **2d: New! Hierarchical Allocation**
    - maybe, if you can find global ordering
- Approach 3: Detection and Recovery
  - Runtime techniques
  - 3a: Ad Hoc Genius
  - **3b: New! Cycle Detection (with recovery)**
    - OK, but then what? need safe rollback mechanism
- Approach 4: Deadlock Avoidance
  - runtime techniques
  - **4a: New! Dijkstra's "Banker's Algorithm"**
    - not generally useful, but OK for specialized applications & OS services

# OK, what do we do?

## Which tools are best?

- Deadlock prevention via hierarchical allocation seems in some sense “best”
  - guarantees that a deadlock can never occur as long as rules are followed
  - no other system support required
- But there’s a significant limitation: all we have is a set of rules for writing *new* programs, not a way to test an arbitrary system for deadlocks
  - in fact, that would be provably impossible!
    - remember the halting problem
    - same reason we don’t have SJF schedulers

# Approach 1 (ignore problem) looking better...

- Most general computing systems don't do deadlock prevention or avoidance as a system service
  - instead they leave applications to fend for themselves
  - avoidance and recovery techniques must be done directly by any applications that require it
- Deadlock is an issue to the OS designer mainly for managing resources *inside* the OS
  - deadlocks within the OS are very disruptive

What about Starvation?

# Starvation != Deadlock

- *Deadlock* is a situation where 2 or more processes are “locked” in a situation where neither *can ever* proceed, no matter what the OS does
  - e.g., circular waiting, etc.
- *Starvation*, on the other hand, is a situation where a process that *could* get service just never *does*
  - e.g., the scheduler never picks it, some other process is always selected ahead of it
- A system can be provably free of deadlock but still be subject to starvation!

# So... deadlock prevention may not prevent starvation

- Mechanisms like semaphores and schedulers don't usually make guarantees about who gets service in what order
- In practice, we avoid starvation in three ways:
  - algorithms that take into account starvation
    - e.g., the dining philosophers solution in Tanenbaum
  - “fair” schedules
    - e.g, based on FCFS RR queues
    - doesn't always work
  - random system behavior that perturbs “bad” cycles
    - no guarantees, but often works in practice