



# Parallel Algorithms

Patrick Cozzi  
University of Pennsylvania  
CIS 565 - Spring 2012

## Announcements

- Presentation topics due 02/07
- Homework 2 due 02/13

## Agenda

- Finish atomic functions from Monday
- Parallel Algorithms
  - Parallel Reduction
  - Scan
  - Stream Compression
  - Summed Area Tables

## Parallel Reduction

- Given an array of numbers, design a parallel algorithm to find the sum.
- Consider:
  - *Arithmetic intensity*: compute to memory access ratio

## Parallel Reduction

- Given an array of numbers, design a parallel algorithm to find:
  - The sum
  - The maximum value
  - The product of values
  - The average value
- How different are these algorithms?

## Parallel Reduction

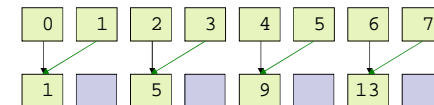
- **Reduction**: An operation that computes a single result from a set of data
- Examples:
  - Minimum/maximum value
  - Average, sum, product, etc.
- **Parallel Reduction**: Do it in parallel. Obviously

## Parallel Reduction

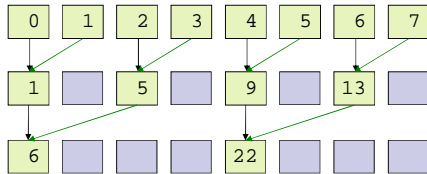
- Example. Find the sum:



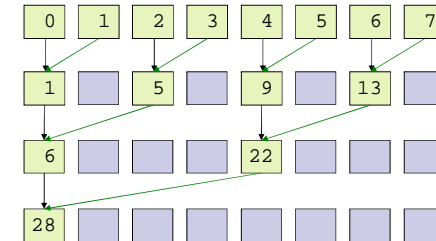
## Parallel Reduction



## Parallel Reduction

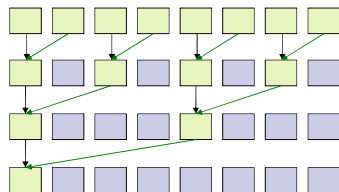


## Parallel Reduction



## Parallel Reduction

- Similar to brackets for a basketball tournament
- $\log(n)$  passes for  $n$  elements



## All-Prefix-Sums

### ■ All-Prefix-Sums

#### □ Input

- Array of  $n$  elements:  $[a_0, a_1, \dots, a_{n-1}]$
- Binary associate operator:  $\oplus$
- Identity:  $I$

#### □ Outputs the array: $[I, a_0, (a_0 \oplus a_1), \dots, (a_0 \oplus a_1 \oplus \dots \oplus a_{n-2})]$

Images from [http://http.developer.nvidia.com/GPUGems3/gpugems3\\_ch39.html](http://http.developer.nvidia.com/GPUGems3/gpugems3_ch39.html)

## All-Prefix-Sums

### ■ Example

- If  $\oplus$  is addition, the array

- [3 1 7 0 4 1 6 3]

- is transformed to

- [0 3 4 11 11 15 16 22]

- Seems sequential, but there is an efficient parallel solution

## Scan

- **Scan**: all-prefix-sums operation on an array of data
- **Exclusive Scan**: Element  $j$  of the result does not include element  $j$  of the input:
  - In: [3 1 7 0 4 1 6 3]
  - Out: [0 3 4 11 11 15 16 22]
- **Inclusive Scan (Prescan)**: All elements including  $j$  are summed
  - In: [3 1 7 0 4 1 6 3]
  - Out: [3 4 11 11 15 16 22 25]

## Scan

- How do you generate an **exclusive scan** from an **inclusive scan**?

- Input: [3 1 7 0 4 1 6 3]

- Inclusive: [3 4 11 11 15 16 22 25]

- Exclusive: [0 3 4 11 11 15 16 22]

- // Shift right, insert identity

- How do you go in the opposite direction?

## Scan

- Use cases

- **Stream compaction**
- **Summed-area tables** for variable width image processing
- **Radix sort**
- ...

## Scan

- Used to convert certain sequential computation into equivalent parallel computation

Sequential	Parallel
<pre> 01. out[0] = 0; 02. for j from 1 to n do 03.   out[j] = out[j-1] + f(in[j-1]); </pre>	<pre> 01. forall j in parallel do 02.   temp[j] = f(in[j]); 03. all_prefix_sums(out, temp); </pre>

Image from [http://http.developer.nvidia.com/GPUGems3/gpugems3\\_ch39.html](http://http.developer.nvidia.com/GPUGems3/gpugems3_ch39.html)

## Scan

- Design a parallel algorithm for exclusive scan
  - In: [3 1 7 0 4 1 6 3]
  - Out: [0 3 4 11 11 15 16 22]
- Consider:
  - Total number of additions

## Scan

- Sequential Scan:** single thread, trivial

```

01. out[0] := 0
02. for k := 1 to n do
03.   out[k] := in[k-1] + out[k-1]

```

- $n$  adds for an array of length  $n$
- Work complexity:**  $O(n)$
- How many adds will our parallel version have?

Image from [http://http.developer.nvidia.com/GPUGems3/gpugems3\\_ch39.html](http://http.developer.nvidia.com/GPUGems3/gpugems3_ch39.html)

## Scan

- Naive Parallel Scan**

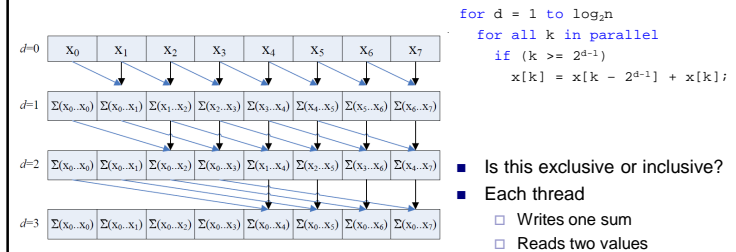


Image from [http://developer.download.nvidia.com/compute/cuda/1\\_1/Website/projects/scan/doc/scan.pdf](http://developer.download.nvidia.com/compute/cuda/1_1/Website/projects/scan/doc/scan.pdf)

## Scan

- *Naive Parallel Scan*: Input



## Scan

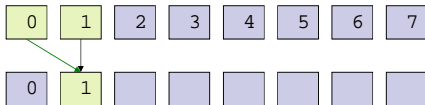
- *Naive Parallel Scan*:  $d = 1, 2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

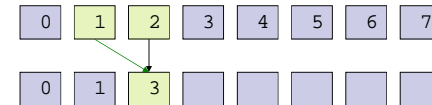
- *Naive Parallel Scan*:  $d = 1, 2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

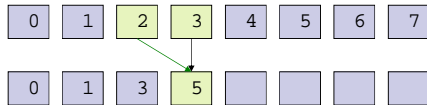
- *Naive Parallel Scan*:  $d = 1, 2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

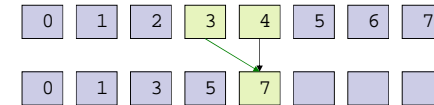
- *Naive Parallel Scan*:  $d = 1$ ,  $2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

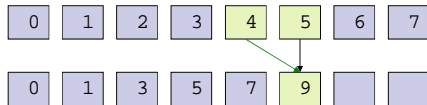
- *Naive Parallel Scan*:  $d = 1$ ,  $2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

- *Naive Parallel Scan*:  $d = 1$ ,  $2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

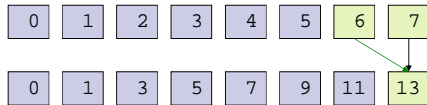
- *Naive Parallel Scan*:  $d = 1$ ,  $2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

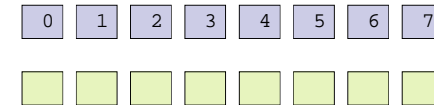
- *Naive Parallel Scan*:  $d = 1$ ,  $2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

- *Naive Parallel Scan*:  $d = 1$ ,  $2^{d-1} = 1$

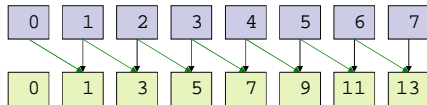


- Recall, it runs in parallel!

```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

- *Naive Parallel Scan*:  $d = 1$ ,  $2^{d-1} = 1$

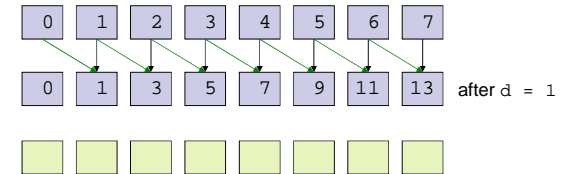


- Recall, it runs in parallel!

```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

- *Naive Parallel Scan*:  $d = 2$ ,  $2^{d-1} = 2$

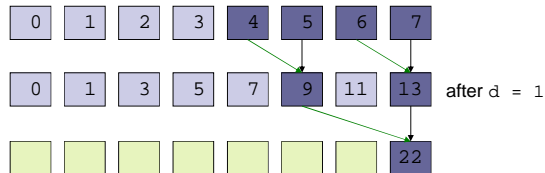


```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```



## Scan

- **Naive Parallel Scan:**  $d = 2, 2^{d-1} = 2$

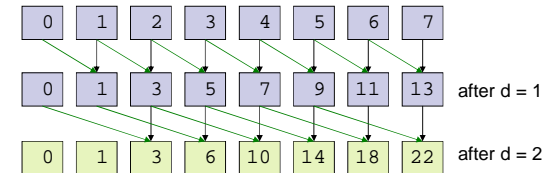


- Consider only  $k = 7$

```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

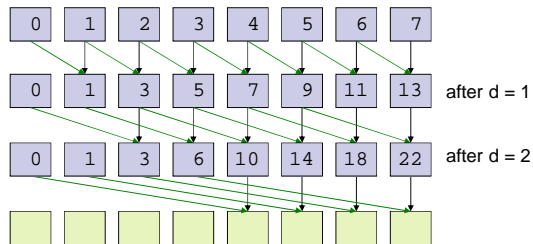
- **Naive Parallel Scan:**  $d = 2, 2^{d-1} = 2$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

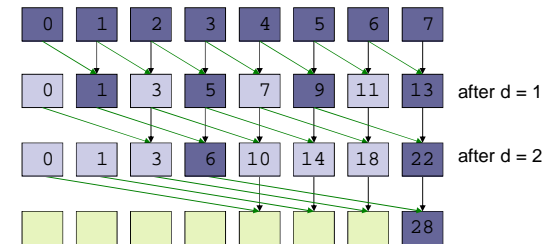
- **Naive Parallel Scan:**  $d = 3, 2^{d-1} = 4$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

- **Naive Parallel Scan:**  $d = 3, 2^{d-1} = 4$

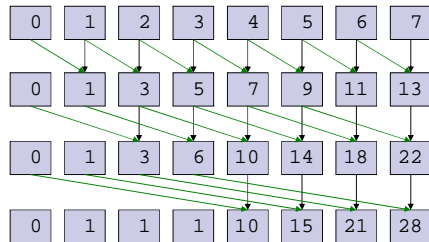


- Consider only  $k = 7$

```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

## Scan

### ■ *Naive Parallel Scan*: Final



## Scan

### ■ *Naive Parallel Scan*

- What is naive about this algorithm?
  - What was the work complexity for sequential scan?
  - What is the work complexity for this?

## Stream Compaction

### ■ *Stream Compaction*

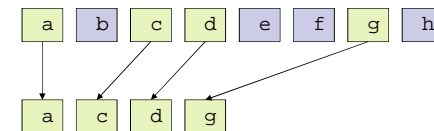
- Given an array of elements
  - Create a new array with elements that meet a certain criteria, e.g. non null
  - Preserve order



## Stream Compaction

### ■ Stream Compaction

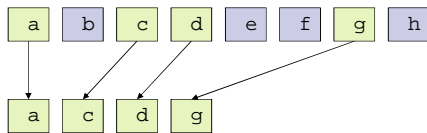
- Given an array of elements
  - Create a new array with elements that meet a certain criteria, e.g. non null
  - Preserve order



## Stream Compaction

### ■ Stream Compaction

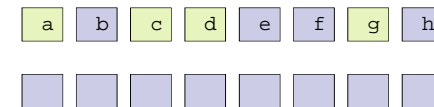
- Used in collision detection, sparse matrix compression, etc.
- Can reduce bandwidth from GPU to CPU



## Stream Compaction

### ■ Stream Compaction

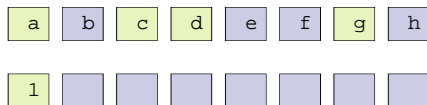
- **Step 1:** Compute temporary array containing
  - 1 if corresponding element meets criteria
  - 0 if element does not meet criteria



## Stream Compaction

### ■ Stream Compaction

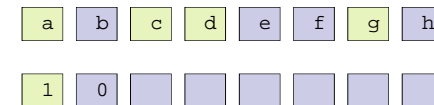
- **Step 1:** Compute temporary array



## Stream Compaction

### ■ Stream Compaction

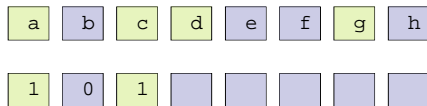
- **Step 1:** Compute temporary array



## Stream Compaction

### ■ Stream Compaction

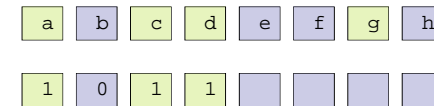
- *Step 1*: Compute temporary array



## Stream Compaction

### ■ Stream Compaction

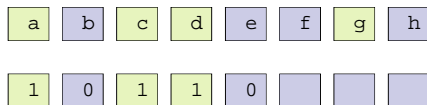
- *Step 1*: Compute temporary array



## Stream Compaction

### ■ Stream Compaction

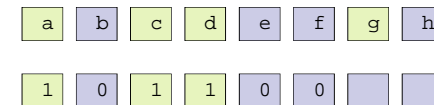
- *Step 1*: Compute temporary array



## Stream Compaction

### ■ Stream Compaction

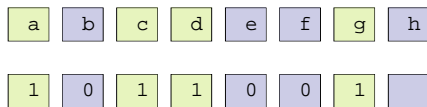
- *Step 1*: Compute temporary array



## Stream Compaction

### ■ Stream Compaction

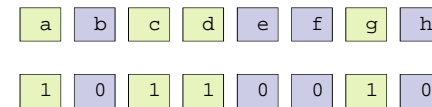
- **Step 1:** Compute temporary array



## Stream Compaction

### ■ Stream Compaction

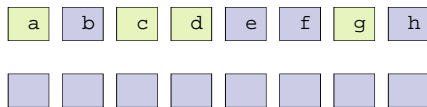
- **Step 1:** Compute temporary array



## Stream Compaction

### ■ Stream Compaction

- **Step 1:** Compute temporary array

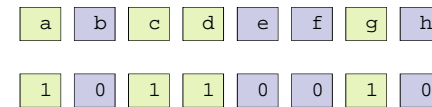


- It runs in parallel!

## Stream Compaction

### ■ Stream Compaction

- **Step 1:** Compute temporary array



- It runs in parallel!

## Stream Compaction

### ■ Stream Compaction

- **Step 2:** Run exclusive scan on temporary array

a	b	c	d	e	f	g	h
1	0	1	1	0	0	1	0
Scan result:							

## Stream Compaction

### ■ Stream Compaction

- **Step 2:** Run exclusive scan on temporary array

a	b	c	d	e	f	g	h
1	0	1	1	0	0	1	0
Scan result:							
0	1	1	2	3	3	3	4

- Scan runs in parallel
- What can we do with the results?

## Stream Compaction

### ■ Stream Compaction

- **Step 3:** Scatter
  - Result of scan is index into final array
  - Only write an element if temporary array has a 1

## Stream Compaction

### ■ Stream Compaction

- **Step 3:** Scatter

a	b	c	d	e	f	g	h
1	0	1	1	0	0	1	0
Scan result:							
0	1	1	2	3	3	3	4
Final array:							
0	1	2	3				

## Stream Compaction

### ■ Stream Compaction

#### □ Step 3: Scatter

	a	b	c	d	e	f	g	h
	1	0	1	1	0	0	1	0
Scan result:	0	1	1	2	3	3	3	4
Final array:	a							
	0	1	2	3				

## Stream Compaction

### ■ Stream Compaction

#### □ Step 3: Scatter

	a	b	c	d	e	f	g	h
	1	0	1	1	0	0	1	0
Scan result:	0	1	1	2	3	3	3	4
Final array:	a	c						
	0	1	2	3				

## Stream Compaction

### ■ Stream Compaction

#### □ Step 3: Scatter

	a	b	c	d	e	f	g	h
	1	0	1	1	0	0	1	0
Scan result:	0	1	1	2	3	3	3	4
Final array:	a	c	d					
	0	1	2	3				

## Stream Compaction

### ■ Stream Compaction

#### □ Step 3: Scatter

	a	b	c	d	e	f	g	h
	1	0	1	1	0	0	1	0
Scan result:	0	1	1	2	3	3	3	4
Final array:	a	c	d	g				
	0	1	2	3				

## Stream Compaction

### ■ Stream Compaction

#### □ Step 3: Scatter

a	b	c	d	e	f	g	h
1	0	1	1	0	0	1	0
Scan result: 0	1	1	2	3	3	3	4

Final array:				
	0	1	2	3

■ Scatter runs in parallel!

## Stream Compaction

### ■ Stream Compaction

#### □ Step 3: Scatter

a	b	c	d	e	f	g	h
1	0	1	1	0	0	1	0
Scan result: 0	1	1	2	3	3	3	4

Final array:	a	c	d	g
	0	1	2	3

■ Scatter runs in parallel!

## Summed Area Table

- Summed Area Table (SAT): 2D table where each element stores the sum of all elements in an input image between the lower left corner and the entry location.

## Summed Area Table

### ■ Example:

Input image					SAT			
<div>2</div>	<div>1</div>	<div>0</div>	<div>0</div>	<div>4</div>	<div>9</div>	<div>12</div>	<div>14</div>	
<div>0</div>	<div>1</div>	<div>2</div>	<div>0</div>	<div>2</div>	<div>6</div>	<div>9</div>	<div>11</div>	
<div>1</div>	<div>2</div>	<div>1</div>	<div>0</div>	<div>2</div>	<div>5</div>	<div>6</div>	<div>8</div>	
<div>1</div>	<div>1</div>	<div>0</div>	<div>2</div>	<div>1</div>	<div>2</div>	<div>2</div>	<div>4</div>	

$$(1 + 1 + 0) + (1 + 2 + 1) + (0 + 1 + 2) = 9$$



## Summed Area Table

### ■ Benefit

- Used to perform different width filters at every pixel in the image in constant time per pixel
- Just sample four pixels in SAT:

$$s_{filter} = \frac{s_{ur} - s_{ul} - s_{lr} + s_{ll}}{w \times h},$$

Image from [http://http.developer.nvidia.com/GPUGems3/gpugems3\\_ch39.html](http://http.developer.nvidia.com/GPUGems3/gpugems3_ch39.html)

## Summed Area Table

### ■ Uses

- Glossy environment reflections and refractions
- Approximate depth of field



Image from [http://http.developer.nvidia.com/GPUGems3/gpugems3\\_ch39.html](http://http.developer.nvidia.com/GPUGems3/gpugems3_ch39.html)

## Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT


## Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

1			

## Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

1	2		

## Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

1	2	2	

## Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

1	2	2	4

## Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

2			
1	2	2	4

## Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

2	5		
1	2	2	4

## Summed Area Table

...

## Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

4	9		
2	6	9	11
2	5	6	8
1	2	2	4

## Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

4	9	12	
2	6	9	11
2	5	6	8
1	2	2	4

## Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

4	9	12	14
2	6	9	11
2	5	6	8
1	2	2	4

## Summed Area Table

How would implement  
this on the GPU?

## Summed Area Table

- Recall *Inclusive Scan*:

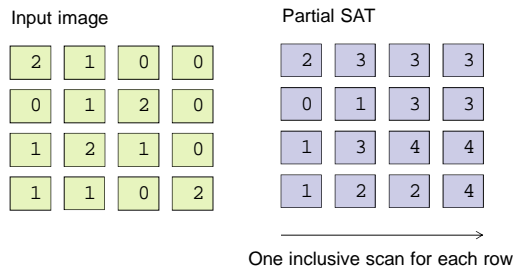
0	1	2	3	4	5	6	7
0	1	3	6	10	15	21	28

## Summed Area Table

How would compute a  
SAT on the GPU using  
inclusive scan?

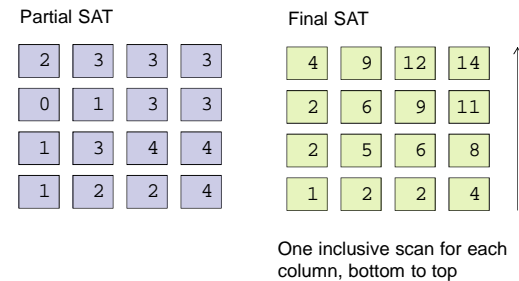
## Summed Area Table

### ■ Step 1 of 2:



## Summed Area Table

### ■ Step 2 of 2:



## Summary

- Parallel reductions and scan are building blocks for many algorithms
- An understanding of parallel programming and GPU architecture yields efficient GPU implementations